InfraSight

Structure of the Code Base, Libraries Used, and Basic Web Development Vocabulary

Technical documentation for project members

1 November 2026

Abstract

InfraSight is a web platform that lets citizens report damaged public infrastructure (potholes, broken street lights, fallen branches) with a photo, and helps municipal staff handle each report from registration to repair. This document explains how the project is built, for a reader who is new to web development. It first defines the basic vocabulary (server, endpoint, API, token, database, and so on). It then describes the overall architecture, the folder structure, the role of every source file and the reason it exists, and the external libraries the project depends on. It ends with a walk-through of one complete user action, and practical notes on running the project and fixing common start-up problems.

Contents

1. Introduction **3**

2. Basic Vocabulary **3**

2.1 Client, server and the network . . . 3

2.2 HTTP: how the client and the server talk . . . 3

2.3 API, endpoint and route . . . 4

2.4 Data and databases . . . 4

2.5 Security . . . 5

2.6 Python project tools . . . 5

2.7 Web server concepts specific to this project . . . 6

2.8 AI and geography . . . 6

3. Overall Architecture **7**

3.1 Layers inside the backend . . . 7

3.2 Life of one request . . . 7

4. Folder Structure **9**

5. Role of Each File **11**

5.1 Root files . . . 11

5.2 Backend core (backend/app/) . . . 11

5.3 Routers (backend/app/routers/) . . . 12

5.4 Services (backend/app/services/) . . . 12

5.5 SAM server (sam_server/) . . . 12

5.6 Frontend (frontend/) . . . 13

6. Libraries Used **14**

6.1 External services . . . 15

7. Walk-through: from a Photo to a Closed Ticket **16**

7.1 A citizen submits a report . . . 16

7.2 Staff process the ticket . . . 16

7.3 Who may do what . . . 16

8. Running the Project and Common Problems **18**

8.1 Starting . . . 18

8.2 When it does not work . . . 18

9. Conclusion **18**

Appendix A. Endpoint Reference **20**

1. Introduction

When a project grows past a few files, it becomes hard to see where to start reading. This document is a map of the InfraSight code. Its goal is that, after reading it, you can open any file and know *what it does*, *who uses it*, and *why it was written that way*.

InfraSight has three users. A **citizen** takes a photo of a problem and places a pin on a map. A **technician** studies the problem, estimates the cost and follows the works. An **elected official** signs the spending decision. Between these people, an artificial intelligence (AI) model reads the photo, proposes a severity, drafts documents and writes updates to the citizen. Every AI output is marked as needing human review.

The document is organised as follows. Section 2 defines the vocabulary used everywhere else; readers who already know web development can skip it. Section 3 gives the architecture. Section 4 shows the folder tree. Section 5 describes every file. Section 6 lists the libraries. Section 7 follows a report through the system. Section 8 explains how to run the project and what to do when it does not start. Appendix A lists all endpoints.

2. Basic Vocabulary

The terms below appear in the code, in error messages and in this document. They are grouped by theme rather than alphabetically, so that each one builds on the previous ones.

2.1 Client, server and the network

Client: the program that asks for something. Here it is the web browser (Chrome, Edge, Firefox) showing the InfraSight console.

Server: the program that answers. Here it is the Python program in `backend/`. It waits for questions, does the work, and sends answers back. A server is a program, not necessarily a separate machine: your own computer can be the client and the server at the same time.

Frontend: the part of the application that runs in the browser and that the user sees: pages, buttons, the map. In this project: the `frontend/` folder.

Backend: the part that runs on the server: it stores data, checks passwords, and calls the AI. The user never sees it directly. In this project: the `backend/` folder.

localhost: the name a computer uses for itself. `http://localhost:8000` means “the server running on this computer”. Nobody else on the internet can reach it.

Port: a number that tells which program on a computer should receive a message, like an apartment number in a building. The InfraSight backend listens on port 8000; the SAM server on port 8100. Only one program can use a given port at a time.

URL: a web address. In `http://localhost:8000/api/tickets/IS-0051`, `http` is the protocol, `localhost` the machine, `8000` the port, and `/api/tickets/IS-0051` the *path*, which says which resource is wanted.

2.2 HTTP: how the client and the server talk

HTTP: HyperText Transfer Protocol, the language of the web. The client sends a *request*; the server sends back exactly one *response*.

Request: a message from the client. It contains a *method*, a *path*, *headers* (extra information such as the login token) and sometimes a *body* (the data being sent, such as a photo).

Response: the server's answer. It contains a *status code*, headers, and a body (usually JSON).

HTTP method: the verb of the request. `GET` reads something without changing it. `POST` creates something or triggers an action. `PUT` replaces or updates something. `DELETE` removes something.

Status code: a three-digit number in the response. `200` OK. `401` not signed in. `403` signed in but not allowed (for example, a citizen trying to sign a decision). `404` not found. `409` conflict (for example, the checklist is not complete). `413` the photo is too large. `422` the data sent is invalid. `429` too many requests. `500` a bug or crash on the server.

JSON: JavaScript Object Notation, a text format for structured data, readable by humans and by every programming language. Example: `{"ref": "IS-0051", "sev": 4}`.

Form data / multipart: another body format, used when a request carries a file. The citizen's photo is sent this way.

2.3 API, endpoint and route

API: Application Programming Interface: the list of questions a program accepts and the shape of its answers. The InfraSight API is the set of URLs under `/api/`. The frontend only talks to the backend through it.

Endpoint: one entry of the API: one method plus one path, for example `POST /api/tickets` (create a report) or `GET /api/stats` (read statistics). Each endpoint is one Python function in the backend. Appendix A lists all of them.

Path parameter: a variable part of a path. In `/api/tickets/{ref}`, `{ref}` is replaced by a real ticket number such as `IS-0051`.

Query parameter: extra values after a `?` in the URL, as in `/api/files/IS-0051.jpg?exp=...&sig=....`

Route: the link between a URL and the function that handles it. The word is often used as a synonym of endpoint.

Router: a group of related routes. InfraSight has four: authentication, tickets, statistics, insights (folder `backend/app/routers/`).

REST: a common style of API design where URLs name things (tickets, zones) and HTTP methods say what to do with them. InfraSight follows this style loosely.

2.4 Data and databases

Database: a program or file that stores data permanently and lets you search it. InfraSight uses SQLite, which keeps the whole database in one file, `backend/infrasight.db`.

Table, row, column: a database is made of tables, like spreadsheet sheets. Each row is one item (one ticket, one user); each column is one property (title, severity).

Primary key: the column that identifies a row uniquely. For tickets it is `ref` (e.g. `IS-0051`).

Foreign key: a column that points to a row of another table. `Ticket.reporter_id` points to the user who reported it.

SQL: Structured Query Language, the language used to talk to most databases.

ORM: Object-Relational Mapper. A library that lets you use database rows as normal Python objects instead of writing SQL by hand. InfraSight uses SQLAlchemy as its ORM.

Session (database): a short conversation with the database during one request. Changes are saved only when the code calls `commit()`.

Migration: a controlled change to the structure of the database (adding a column, for example). The project has a temporary helper for this, `add_missing_columns()`, until a real migration tool (Alembic) is added.

2.5 Security

Authentication: checking *who* you are (username and password).

Authorization: checking *what you are allowed to do* once known (a citizen may report; only staff may advance a ticket; only an elected official may sign the decision stage).

Role: the category of a user: `citizen`, `tech` (technician) or `elected`. Technicians and elected officials together are called *staff*.

Password hash: a one-way transformation of a password. The database stores the hash, never the password, so a stolen database does not reveal passwords. InfraSight uses the *scrypt* algorithm with a random *salt* (a random value mixed in so that two identical passwords give different hashes).

Token / JWT: after a successful login, the server gives the browser a *token*: a signed piece of text that proves who the user is. The browser sends it with every request in the header `Authorization: Bearer <token>`. InfraSight uses JSON Web Tokens (JWT) that expire after 12 hours.

Secret key: a long random value known only by the server, used to sign tokens and photo links. If it changes, every user is logged out. It lives in `backend/.env`.

Signed URL: a link that contains its own proof of permission and an expiry date. Used for photos, because an HTML `<img>` tag cannot send a token.

CORS: Cross-Origin Resource Sharing. Browsers block a page from calling a server on a different address unless that server allows it. The backend lists the allowed addresses in `cors_origins`.

Rate limit: a maximum number of requests per user and per period. InfraSight allows 30 AI calls per user per hour, because AI calls cost money.

API key: a secret password that identifies the project to an external paid service (here the AI provider). It must never be committed to Git or shown in the browser.

2.6 Python project tools

Library / package / dependency: code written by other people that the project reuses. A project's *dependencies* are the libraries it needs to run. They are listed in `requirements.txt`.

pip: the Python tool that downloads and installs libraries: `pip install -r requirements.txt`.

Virtual environment (venv): a private folder of installed libraries for one project, so that projects do not conflict. Here: `backend/.venv/`. It is not committed to Git; `start.bat` recreates it.

Environment variable / .env file: a setting given to a program from outside its code: API keys, the secret key, the database location. The file `.env` holds the real values and stays private; `.env.example` is a template with empty values that is committed so others know what to fill in.

Module: one Python file. A *package* is a folder of modules with an `__init__.py` file (often empty) that tells Python the folder can be imported.

Import: the instruction that lets one file use code from another, e.g. `from .db import Ticket`. The dot means "from the same package".

2.7 Web server concepts specific to this project

Framework: a library that gives a ready-made structure for an application, so you only write the parts that are specific to it. FastAPI is the backend framework.

ASGI server: the program that actually listens on the port and passes requests to the framework. Here: Uvicorn.

Static files: files sent to the browser as they are, without computation: HTML, CSS, JavaScript, images. The backend serves the `frontend/` folder as static files.

Middleware: code that runs on every request before or after the endpoint. InfraSight uses one: the CORS middleware.

Dependency injection: in FastAPI, a function parameter such as `u: User = Depends(staff)` means "before running this endpoint, run `staff()` and give me its result". It is how the project checks the token and the role in one line.

Schema / validation: a description of the expected shape of data (field names, types, allowed values). Pydantic checks incoming data against it and rejects bad data with a 422 error. It also checks that the AI answered in the expected format.

Background task: work started after the response is sent, so the user does not wait. Citizen notifications are written this way.

Thread pool: a set of helper threads used to run slow, blocking work (network calls to the AI or to OpenStreetMap) without freezing the server for other users.

2.8 AI and geography

AI model / LLM: a large language model that reads text and images and writes text. InfraSight can use Meta's Muse Spark or Anthropic's Claude. Note that the word *model* also means a database table class in the code; the context tells which.

Prompt: the instruction text sent to the AI. All prompts are kept on the server in `services/ai.py`, never in the browser.

Triage: the first quick classification of a report: type, severity (1 to 5), urgency, who owns the land.

Segmentation: finding the exact outline of an object in a photo. SAM (Segment Anything Model) does this; it draws the outline of the pothole.

GPU: a graphics card. Neural networks such as SAM run much faster on it, which is why SAM runs on a separate machine.

Geocoding: turning a map position (latitude, longitude) into a street address, or the reverse.

Latitude / longitude: the two numbers that give a position on Earth (`lat`, `lng`).

3. Overall Architecture

InfraSight is split into three programs plus external services. Figure 1 shows them and how they communicate.

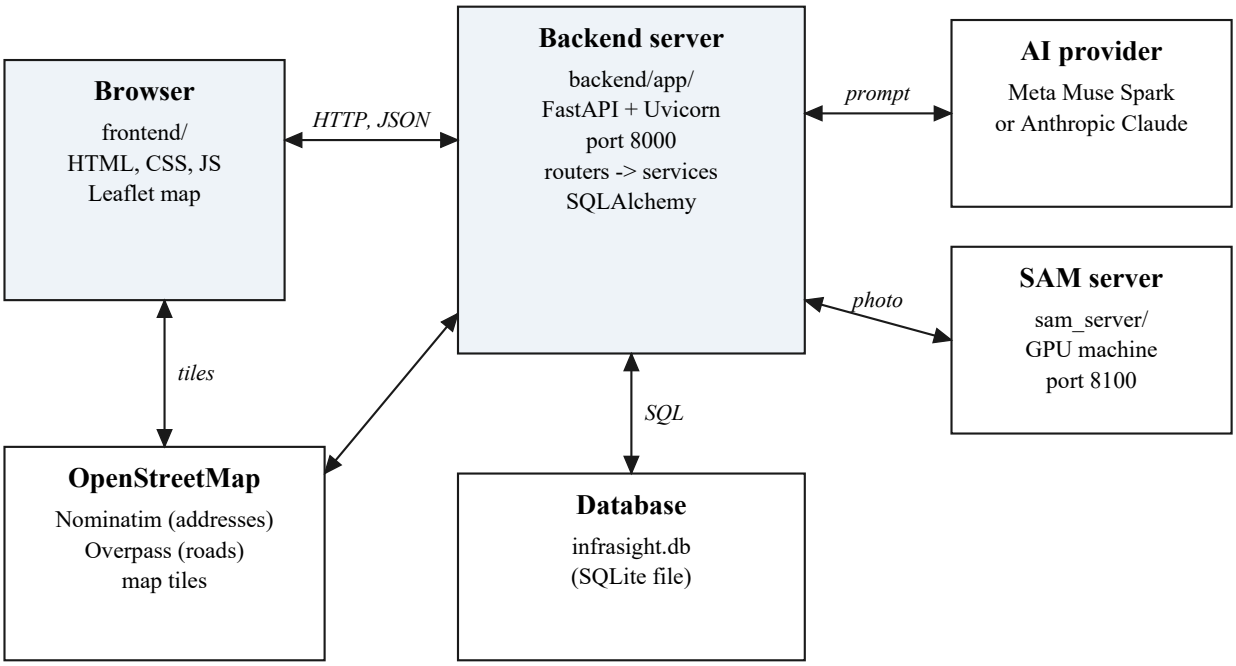


Figure 1. Components of InfraSight. Arrows are network calls; the backend is the only component that talks to all the others.

The browser loads the frontend files and displays the console. It never talks to the AI or to the database directly. It only calls the backend's API.

The backend is the centre. It checks who is calling, reads and writes the database, and calls the external services. It also serves the frontend files, so a single address, `http://localhost:8000`, is enough to use the whole application.

The SAM server is optional. It needs a powerful graphics card, so it is designed to run on another machine. If it is not configured, the application still works; it simply does not draw damage outlines.

This separation has three benefits. *Security*: secrets (API keys, password hashes) stay on the server. *Cost control*: the server can limit AI calls per user. *Replaceability*: the AI provider can be changed in one setting without touching the frontend.

3.1 Layers inside the backend

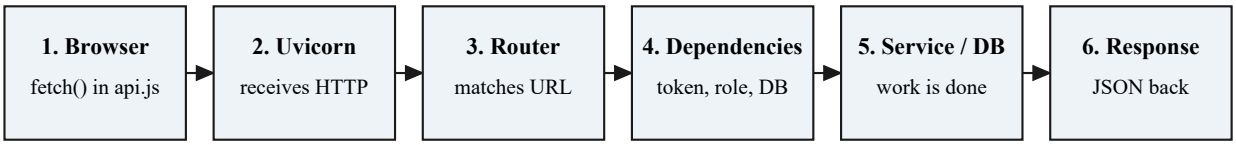
Inside `backend/app/`, code is organised in layers. Each layer only calls the layers below it:

- **Routers** (`routers/`): receive HTTP requests, check permissions, and decide what to do.
- **Services** (`services/`): do one technical job each (call the AI, find an address, compute a detour). They know nothing about HTTP users or roles.
- **Core modules**: database tables (`db.py`), security (`auth.py`), settings (`config.py`), data shapes (`schemas.py`), and rules (`regulations.py`).

This makes each file small and focused: to change how addresses are found, you only open `geo.py`.

3.2 Life of one request

Figure 2 shows what happens when a technician clicks “Advance to next stage”.




```
Example: POST /api/tickets/IS-0051/advance  with header  Authorization: Bearer <token>
Answer: 200 OK and the updated ticket as JSON, or 401 / 403 / 409 with an error message
```

Figure 2. Path of one HTTP request through the backend.

The JavaScript function in `api.js` sends the request with the token. Uvicorn receives it and gives it to FastAPI, which finds the matching function in `routers/tickets.py`. Before that function runs, its dependencies run: `staff()` decodes the token and refuses citizens; `get_db()` opens a database session. The function then checks the business rules (checklist complete, elected official for the decision stage), saves the change, schedules a notification in the background, and returns the updated ticket as JSON.

4. Folder Structure

Figure 3 shows the project tree, limited to the parts currently in use. Files ignored by Git (not shared) are marked.

```
InfraSight/
├── start.bat                 one-click launcher for Windows
├── README.md                 short presentation and run instructions
├── .gitignore                files Git must not track (secrets, database, venv)
├── InfraSight-pitch.pdf      presentation deck
│
├── frontend/                 what runs in the browser (no build step)
│   ├── index.html            the single page; loads the CSS and scripts
│   ├── css/styles.css        colours, layout, typography
│   └── js/
│       ├── api.js            talks to the backend (one function per endpoint)
│       ├── app.js            screens, views, charts, login
│       ├── mapview.js        the main city map (Leaflet)
│       └── map.js            drawing of danger zones on maps
│
├── backend/                  the server (Python, FastAPI)
│   ├── requirements.txt      libraries to install
│   ├── .env.example          template of settings (committed)
│   ├── .env                  real settings and keys     [ignored by Git]
│   ├── .venv/                installed libraries        [ignored by Git]
│   ├── infrasight.db         the SQLite database        [ignored by Git]
│   ├── uploads/              citizens' photos           [ignored by Git]
│   └── app/
│       ├── main.py           entry point: builds the application
│       ├── config.py         reads settings from .env
│       ├── db.py             database tables
│       ├── auth.py           passwords, tokens, roles, rate limit
│       ├── schemas.py        expected shape of AI answers
│       ├── serializers.py    database row -> JSON for the browser
│       ├── regulations.py    the 7 stages, legal notes, directory
│       ├── seed.py           demo data
│       ├── cli.py            command-line account management
│       ├── routers/          the API endpoints
│       │   ├── auth.py       login, register, profile
│       │   ├── tickets.py    everything about a report
│       │   ├── stats.py      dashboard numbers
│       │   └── insights.py   hotspots, weekly brief, notifications
│       └── services/         technical helpers
│           ├── ai.py         all AI prompts and calls
│           ├── vision.py     client for the SAM server
│           ├── geo.py        addresses and street shapes
│           ├── routing.py    detour computation
│           ├── images.py     photo checking and resizing
│           ├── notify.py     updates to citizens
│           └── export.py     printable HTML report
│
├── sam_server/               optional image segmentation server (GPU)
│   ├── server.py
│   └── requirements.txt
│
└── client/, server/, docs/   earlier version of the project (not used by start.bat)
```

Figure 3. Project tree. Indentation shows which folder contains which file.

Why separate folders? The frontend and the backend are written in different languages (JavaScript and Python), run in different places (browser and server), and are installed differently. Keeping them

apart makes it clear which code runs where. The SAM server is separate again because it needs heavy libraries (PyTorch, about two gigabytes) and a GPU that the main server does not need.

Why are some files ignored? `.env` contains secrets. `infrasight.db` and `uploads/` contain personal data and change all the time. `.venv/` is large and can be rebuilt from `requirements.txt`. None of these should be in a shared repository.

Note on older folders. `client/`, `server/` and `docs/` belong to an earlier version of the project and are still in the repository. The current application, started by `start.bat`, uses only `frontend/`, `backend/` and optionally `sam_server/`.

5. Role of Each File

5.1 Root files

`start.bat`: a Windows script. On the first run it creates the virtual environment, installs the libraries, and creates `backend/.env` with a random secret key and the demo password `demo1234`. On every run it starts Uvicorn on port 8000. *Why:* so that a non-developer can start the project with a double-click.

`README.md`: the first page shown on GitHub: the folder layout and how to run the project on other systems.

`.gitignore`: the list of files Git must not track. *Why:* to keep secrets and personal data out of the repository.

5.2 Backend core (backend/app/)

`main.py`: the entry point that Uvicorn loads (`app.main:app`). At start-up it creates the database tables if they do not exist, adds missing columns, inserts demo data if asked, and starts a background thread that fills in missing street addresses. It then creates the FastAPI application, adds the CORS middleware, attaches the four routers, defines two small endpoints (`/api/health`, `/api/meta`) and finally serves the `frontend/` folder at the root address. *Why last:* so that `/api/...` URLs are matched before the static files.

`config.py`: defines a `Settings` class listing every setting with a default value, and reads the real values from `.env`. Every other file imports `settings` from here. *Why:* one place for configuration, with types checked automatically.

`db.py`: describes the database tables as Python classes: `User`, `Ticket`, `Audit` (history of every action on a ticket), `Notification`, `Insight` (saved AI outputs such as the weekly brief) and `Zone` (danger zones drawn on the map). It also provides `get_db()`, which opens a database session for each request. A ticket stores its main fields in columns and less structured AI results in a flexible JSON column called `data`.

`auth.py`: security functions: hashing and checking passwords (scrypt), creating and reading tokens (JWT), the dependencies `current_user` (refuses unsigned requests) and `staff` (refuses citizens), signed photo links, and the AI rate limit (30 calls per hour per user, kept in memory).

`schemas.py`: Pydantic classes describing the exact JSON the AI must return: `Triage`, `Report`, `Consequences`, `DangerZone`, `SizeEstimate`, `RepairCheck`, `Forecast`. *Why:* an AI can answer in an unexpected format; validating it protects the rest of the code.

`serializers.py`: converts a database `Ticket` into the JSON shape the frontend expects, and **filters it according to who is asking**. Staff and the reporting citizen see everything; other citizens see only public fields (no photo, no history, no AI output). It also builds the context given to the AI.

`regulations.py`: the business rules in one place: the seven stages of a repair (intake, study, decision, procurement, authorizations, execution, acceptance), the checklist for each stage, the legal notes (French public procurement thresholds, local government code), the directory of responsible authorities, and the land owner categories. The frontend gets this through `GET /api/meta`, except the AI instructions, which stay on the server.

`seed.py`: inserts sample tickets and, if `DEMO_PASSWORD` is set, three demo accounts (`citizen`, `tech`, `elected`). Runs only on an empty database with `SEED_DEMO` enabled. Never used in production.

`cli.py`: a command-line tool to create staff accounts and change roles or passwords, e.g. `python -m app.cli create-user alice tech "Alice"`. *Why:* the website only lets people register as citizens; staff accounts must be created by an administrator.

`__init__.py`: empty files that mark `app/`, `routers/` and `services/` as Python packages.

5.3 Routers (backend/app/routers/)

`auth.py`: login, registration (always as a citizen), reading the current user, and changing the language of notifications (nine languages including French, Breton and Arabic).

`tickets.py`: the largest file. Creating a report (photo, address lookup, AI triage, duplicate detection); reading tickets; the full AI report; damage outline and size; after-repair photo check; detour; AI help for each stage; checklist; advancing a stage; exporting a report; saving danger zones; and serving photos through signed links. Each AI-written log line is tagged "AI-generated, human review required".

`stats.py`: one endpoint that counts tickets by type, stage and severity for the staff dashboard.

`insights.py`: city-wide views for staff: hotspots (streets with many problems) and an AI forecast of where problems will appear next, the weekly brief for elected officials, and the citizen's notification inbox.

5.4 Services (backend/app/services/)

`ai.py`: every prompt in the project and the code that sends it to the AI provider (Meta or Anthropic, chosen by `AI_PROVIDER`). Functions: `triage`, `full_report`, `stage_assist`, `estimate_size`, `check_repair`, `traffic_order`, `forecast`, `citizen_update`, `weekly_brief`. It asks for JSON, then validates the answer with the schemas. *Why one file:* prompts are easy to find, review and improve.

`vision.py`: sends the photo to the SAM server with a short word chosen from the ticket type ("pothole", "street lamp") and returns the outline. Does nothing if `SAM_URL` is empty.

`geo.py`: uses OpenStreetMap's Nominatim service to turn a pin into an address and to get the shape of the street, so a danger zone can follow the street. Respects the service's limit of one request per second.

`routing.py`: computes a detour around a blocked street: downloads nearby roads from OpenStreetMap (Overpass), removes the blocked part, and finds the shortest path with Dijkstra's algorithm.

`images.py`: checks the uploaded photo (size limit, real image) and shrinks it to a 1280-pixel JPEG, stored once.

`notify.py`: writes a short update in the citizen's language when something happens to their report. Runs as a background task.

`export.py`: builds a printable HTML report of a ticket, with the photo embedded.

5.5 SAM server (sam_server/)

`server.py`: a second, small FastAPI application. It loads the SAM 3.1 model once at start-up, then answers `POST /segment`: given a photo and a word, it returns up to five outlines as polygons, with coordinates between 0 and 1 so they do not depend on the photo's size. OpenCV simplifies each outline to fewer points. *Why separate:* it needs a GPU and very large libraries.

5.6 Frontend (frontend/)

`index.html`: the only HTML page. It creates an empty layout (side menu and main area) and loads the fonts, Leaflet, the stylesheet and the four scripts in order. All content is then drawn by JavaScript. This is called a *single-page application*.

`css/styles.css`: the visual design: colours, spacing, the five-step severity colour scale.

`js/api.js`: the only file that talks to the server. It stores the login token in the browser, adds it to every request, turns server errors into readable messages, and exposes one function per endpoint under the global object `CF` (e.g. `CF.advance(ref)`). If the server answers 401, it logs out and reloads.

`js/app.js`: the application itself: login screen, navigation, ticket list, ticket detail tabs, stage checklists, charts drawn as SVG, notifications. It calls `api.js` and the map scripts.

`js/mapview.js`: the main city map built with Leaflet: markers grouped into clusters, a dark style for staff, filters, and the pin a citizen places to report a problem.

`js/map.js`: draws danger zones as irregular coloured shapes (yellow at the edge to deep red at the centre) and the small map in a ticket's “Danger” tab. Shapes are generated from the position, so a zone always looks the same.

6. Libraries Used

A library is chosen when writing the same code ourselves would be long, error-prone, or less secure. Table 1 lists the backend libraries from `backend/requirements.txt`, Table 2 those of the SAM server, and Table 3 those loaded by the browser.

Table 1. Backend libraries.

Library	What it is	Where and why it is used
`fastapi`	Web framework for building APIs in Python.	All routers and `main.py`. Chosen because it is fast, checks input data automatically from type hints, and generates interactive API documentation at `/docs`.
`uvicorn[standard]`	ASGI server.	Listens on port 8000 and passes requests to FastAPI. Started by `start.bat`.
`python-multipart`	Parser for form data.	Needed by FastAPI to receive the photo upload in `POST /api/tickets`.
`sqlalchemy`	ORM and database toolkit.	`db.py` and every router. Lets the code use Python objects instead of SQL, and makes a later move from SQLite to PostgreSQL a one-line change.
`pydantic-settings`	Settings loader built on Pydantic.	`config.py`: reads `.env` and checks types (a number stays a number).
`pyjwt`	JSON Web Token library.	`auth.py`: creates and verifies login tokens.
`pillow`	Image processing (PIL).	`images.py`: opens, validates, resizes and recompresses photos.
`httpx`	HTTP client.	`vision.py`: sends the photo to the SAM server.
`openai`	Client for OpenAI-compatible APIs.	`ai.py`: used to call Meta's Model API, which follows the same format.
`anthropic`	Official Claude client.	`ai.py`: used when `AI_PROVIDER=anthropic`.

Pydantic itself (data validation) is installed automatically with FastAPI. Some features use only Python's standard library, which needs no installation: `hashlib` (scrypt password hashing), `urllib` (OpenStreetMap calls), `heapq` (shortest path), `threading`, `json`.

Table 2. SAM server libraries.

Library	Where and why it is used
`fastapi, uvicorn, python-multipart`	Same roles as in the backend: the SAM server is a small API.
`torch (PyTorch)`	Runs the neural network on the GPU (or the CPU, much slower).
`transformers, accelerate`	Hugging Face libraries that download the SAM 3.1 model and prepare its inputs and outputs.
`pillow`	Opens the uploaded photo.
`numpy`	Fast arrays of numbers; the mask returned by SAM is converted to one.
`opencv-python-headless`	Finds the contour of the mask and simplifies it to a polygon. “Headless” means without a graphical window, which a server does not need.

Table 3. Frontend libraries, loaded from public CDNs in index.html.

Library	Where and why it is used
Leaflet 1.9.4	Interactive maps (zoom, markers, shapes) in `mapview.js` and `map.js`. Free, light, and works with OpenStreetMap tiles.
Leaflet.markercluster 1.5.3	Groups nearby markers into one numbered circle when zoomed out, so the map stays readable with many tickets.
Google Fonts (Inter, Newsreader)	Typefaces for the interface.

The frontend uses no framework (no React, no build step). *Why:* the files can be edited and reloaded directly, and there is nothing to compile or install for the browser side.

6.1 External services

- **Meta Model API (Muse Spark)** or **Anthropic API (Claude)**: the AI models. Paid, need an API key in `.env`.
- **OpenStreetMap Nominatim**: addresses from map positions. Free, limited to one request per second.
- **OpenStreetMap Overpass**: road network data for detours. Free; two servers are tried in order.
- **Hugging Face**: hosts the SAM 3.1 model, downloaded once by the SAM server.

7. Walk-through: from a Photo to a Closed Ticket

7.1 A citizen submits a report

This is the most complete action in the project; following it touches almost every file.

- In the browser, the citizen places a pin, writes a description and chooses a photo. `app.js` collects them in a form and calls `CF.create()` in `api.js`, which sends `POST /api/tickets` with the token.
- FastAPI calls `create()` in `routers/tickets.py`. The dependency `ai_quota` checks the token and the hourly AI limit.
- `images.normalize()` checks and shrinks the photo. `geo.place()` finds the address of the pin.
- The open tickets are sorted by distance. `ai.triage()` sends the photo, the description and the list of nearby tickets to the AI, which returns a title, type, severity, urgency, land owner, consequences, a danger zone and possibly the number of a duplicate. `schemas.Triage` validates the answer.
- If the AI names a duplicate that is open and within 80 metres, the report is merged into it and the citizen becomes a follower. Otherwise a new ticket `IS-00NN` is created, the photo is saved in `uploads/`, and history lines are added to `audit_log`.
- The response is sent at once. Afterwards, a background task in `notify.py` writes a confirmation in the citizen's language.
- If the AI is unavailable, the report is still saved, with default values, for manual triage. *Why:* a citizen's report must never be lost because of an external service.

7.2 Staff process the ticket

The ticket then moves through the seven stages defined in `regulations.py`. At each stage, staff can ask the AI for help (`/stages/{stage}/assist`), which drafts the documents of that stage. To move on (`/advance`), every checklist item must be ticked, and the decision stage (stage 3, arbitration) can only be signed by an elected official. When the last stage is complete the ticket is closed and the citizen is informed.

Table 4. The seven stages of a repair (source: regulations.py).

#	Stage	Typical output
1	Intake and registration	Ticket number, routing, duplicate check, acknowledgement
2	Technical and financial study	Inspection, cost estimate, land ownership check
3	Arbitration and decision	Purchase order or council deliberation; elected sign-off
4	Public procurement	Technical specification (CCTP), choice of procedure, contract
5	Authorizations	Network operator declarations (DT-DICT), traffic order
6	Execution and monitoring	Service order, site visits, photos
7	Acceptance and closure	Acceptance report, payment, message to the citizen

7.3 Who may do what

Table 5. Permissions by role (enforced in auth.py and the routers, not only in the interface).

Action	Citizen	Technician	Elected
Register, sign in, report a problem	yes	yes	yes
See own reports in full, receive updates	yes	yes	yes
See other people's photos and history	no	yes	yes
Run AI report, damage outline, detour, stage help	no	yes	yes
Tick checklists, advance a stage	no	yes	yes
Sign the decision stage	no	no	yes
Statistics, hotspots, weekly brief	no	yes	yes

Permissions are checked on the server. Hiding a button in the browser is only a convenience: a user could still send the request by hand, so the server must refuse it.

8. Running the Project and Common Problems

8.1 Starting

On Windows, double-click `start.bat`. Leave the black window open: it *is* the server, and closing it stops the site. When it shows `Uvicorn running on http://127.0.0.1:8000`, open `http://localhost:8000` in a browser and sign in with `citizen`, `tech` or `elected` and the `DEMO_PASSWORD` from `backend/.env`. To use the AI features, put a real `META_API_KEY` (or `ANTHROPIC_API_KEY` with `AI_PROVIDER=anthropic`) in `backend/.env` and restart.

The interactive API documentation, generated by FastAPI, is at `http://localhost:8000/docs`. It lists every endpoint and lets you try them, which is a good way to learn the API.

8.2 When it does not work

Table 6. Common start-up problems.

Symptom	Probable cause	What to do
The page is blank or shows errors after opening `frontend/index.html` directly	The page was opened as a file, not through the server.	Open `http://localhost:8000` instead.
"The InfraSight server cannot be reached"	The server is not running, or its window was closed.	Run `start.bat` again and keep the window open.
Error `address already in use` / port 8000	Another copy of the server (or another program) already uses port 8000.	Close the other window, or find it with `netstat -ano | findstr :8000` and end that process. If the old copy is InfraSight, just use it.
"Setup failed. Is Python 3.10+ installed?"	Python is missing, or a library failed to install.	Install Python 3.12 and tick "Add to PATH". Very new versions (3.14) may lack ready-made packages for some libraries; 3.12 is the safest.
Login refused	Demo users are created only with `SEED_DEMO=true` and a `DEMO_PASSWORD`, on an empty database.	Check `backend/.env`, or create a user with `python -m app.cli`.
Report saved "for manual triage"	No AI key, wrong key, or AI limit reached.	Put a valid key in `.env` and restart.
Everybody is logged out after a restart	`SECRET_KEY` is empty, so a new random one is made each time.	Set `SECRET_KEY` in `.env` (`start.bat` does it on first run).
No damage outline	`SAM_URL` is not set.	Optional: run the SAM server and set `SAM_URL`.

9. Conclusion

InfraSight follows a classic and robust structure: a static frontend that only talks to an API, a backend divided into routers, services and core modules, and a separate optional server for heavy image processing. Rules and prompts are each kept in a single file (`regulations.py` and `services/ai.py`), security decisions are taken on the server, and AI output is always validated and labelled for human

review.

Known limits, marked as TODO in the code, are: no real database migration tool yet (Alembic), SQLite instead of PostgreSQL for production, a rate limit kept in memory (one server process only), and legal figures that still need validation by a legal department. A reader who wants to go further should start by opening `http://localhost:8000/docs`, then read `routers/tickets.py` alongside Section 7 of this document.

Appendix A. Endpoint Reference

All paths start with `http://localhost:8000`. “Public” needs no login; “user” needs any signed-in account; “staff” needs a technician or elected official; “+AI” counts toward the hourly AI limit.

Table 7. All backend endpoints.

Method	Path	Access	Purpose
POST	/api/auth/login	public	Sign in; returns a token and the user
POST	/api/auth/register	public	Create a citizen account
GET	/api/auth/me	user	Current user
PUT	/api/auth/me	user	Change notification language
GET	/api/health	public	Is the server alive?
GET	/api/meta	public	Stages, checklists, directory, owners, languages
GET	/api/state	user	All tickets visible to this user, and zones
GET	/api/tickets/{ref}	user	One ticket
POST	/api/tickets	user +AI	New report (photo, pin, description); AI triage
POST	/api/tickets/{ref}/merge	staff	Merge a ticket into another (duplicate)
GET	/api/files/{name}	signed link	A photo (needs exp and sig parameters)
POST	/api/tickets/{ref}/report	staff +AI	Full AI report
POST	/api/tickets/{ref}/damage	staff +AI	Damage outline (SAM) and size estimate
POST	/api/tickets/{ref}/after-photo	staff +AI	After-repair photo; AI compares before and after
POST	/api/tickets/{ref}/detour	staff +AI	Detour route and draft traffic order
POST	/api/tickets/{ref}/stages/{stage}/assist	staff +AI	AI help for one stage
PUT	/api/tickets/{ref}/checklist	staff	Tick or untick a checklist item
POST	/api/tickets/{ref}/advance	staff	Move to next stage (elected for decision)
POST	/api/tickets/{ref}/apply-severity	staff	Apply the severity proposed by the AI report
GET	/api/tickets/{ref}/report.html	staff	Printable report
PUT	/api/zones	staff	Save danger zones drawn on the map
GET	/api/stats	staff	Dashboard counts
GET	/api/hotspots	staff	Streets with many problems, and last forecast
POST	/api/hotspots/forecast	staff +AI	New AI forecast
GET	/api/brief	staff	Last weekly brief
POST	/api/brief	staff +AI	Generate a new weekly brief
GET	/api/notifications	user	My updates
POST	/api/notifications/read	user	Mark my updates as read

SAM server (port 8100, separate machine): `GET /health` and `POST /segment` (photo + word → outlines).